

## Exercise Sheet 4

---

Mario Rodríguez Ruiz

May 11, 2017

## Contents

<b>1</b>	<b>Explain your system</b>	<b>3</b>
<b>2</b>	<b>Calculating Pi using collective communication</b>	<b>3</b>
2.1	How you assign tasks to different processes? . . . . .	3
2.2	How you combine results from all the processes? . . . . .	4
2.3	Provide runtime analysis on varying number of processes . . . . .	5
2.3.1	The value of $\infty$ set as $10^2$ . . . . .	5
2.3.2	The value of $\infty$ set as $10^3$ . . . . .	6
2.3.3	Run program . . . . .	6
<b>3</b>	<b>Matrix - matrix multiplication using collective communication</b>	<b>7</b>
3.1	How you assign tasks to different processes? . . . . .	7
3.2	How you combine results from all the processes? . . . . .	8
3.3	Provide runtime analysis on varying number of processes . . . . .	8
3.3.1	First run . . . . .	8
3.3.2	Second run . . . . .	9
3.3.3	Third run . . . . .	10
3.3.4	Conclusions . . . . .	10

## List of Figures

2.1	AllReduce Example . . . . .	4
2.2	Combine results for the value of $\infty$ set as $10^2$ . . . . .	5
2.3	Combine results for the value of $\infty$ set as $10^3$ . . . . .	6
3.1	Broadcast of second matrix . . . . .	7
3.2	Scatter of first matrix . . . . .	7
3.3	Gather of result matrix . . . . .	8
3.4	Combine first run results . . . . .	9
3.5	Combine second run results . . . . .	9
3.6	Combine third run results . . . . .	10

## List of Tables

1.1	My system . . . . .	3
2.1	Results for processes-time for the value of $\infty$ set as $10^2$ . . . . .	5
2.2	Results for processes-time for the value of $\infty$ set as $10^3$ . . . . .	6
3.1	First run results time per processes . . . . .	8
3.2	Second run results time per processes . . . . .	9
3.3	Third run results time per processes . . . . .	10

# 1 Explain your system

Home System	
Machine	Asus Notebook ROG G60Jx
Operating System	Windows 10 Pro 64-bit
CPU	Intel Core i7 720QM @1.60GHz
Number of cores	4
Number of threads	8
RAM	16GB @665MHz (9-9-9-24)
Programming language version Python	v3.6.1:69c0db5 64 bit
Programming language version Java	v1.8

Table 1.1: My system

## 2 Calculating Pi using collective communication

### 2.1 How you assign tasks to different processes?

Each process calculates its own calculation limits depending on the total calculation number and the number of processes that exist.

```
1 # Calculate calculation limits for each process
2 def limits():
3     # Number of calculations for each process
4     n_calc = int(infinity/size)
5     # Remainder of the division of the calculations of each
6     # process
7     rest = int(infinity % size)
8
9     # If there is a remainder, it is divided into
10    # one unit per process until do all allotted.
11    if rest != 0:
12        # If rank is smaller than rest, it is that
13        # part of the rest still has to be distributed.
14        if rank < rest :
15            # Whenever an additional unit is assigned to a process ,
16            # both the
17            # start position and the end position of the other
18            # processes will be affected.
19            n_calc = n_calc + 1
20            if rank == 0:
21                ini = 0
22            else:
23                ini = rank*n_calc
24        # At this moment all the remainder has already been
25        # distributed.
26    else:
27        ini = rank*n_calc+rest
```

```

24     else:
25         ini = rank*n_calc
26         end = ini + n_calc
27         return ini,end

```

Once each process has calculated its calculation limits, they make calls to the function from the beginning of the range (**i**) to the end of it (**end**).

```

1  # Bailey-Borwein-Plouffe formula
2  def bbp(i):
3      eightI = 8*i
4      sixteen = (Decimal(1)/(16**i))
5      one = (Decimal(4)/(eightI+1))
6      four = (Decimal(2)/(eightI+4))
7      five = (Decimal(1)/(eightI+5))
8      six = (Decimal(1)/(eightI+6))
9      return sixteen*(one-four-five-six)
10
11 def calc_pi(i , end):
12     pi = Decimal(0)
13     for i in range(i, end):
14         pi += bbp(i)
15     return pi

```

## 2.2 How you combine results from all the processes?

To combine all the results, I've used **AllReduce** (Figure 2.1). Reduces values on all processes to a single value onto all processes.

### MPI\_Allreduce




Figure 2.1: AllReduce Example

```

1 # Calculation limits for each process
2 ini, end = limits()
3
4 # User alterable precision (infinity)
5 getcontext().prec = (int(infinity))
6 pi = calc_pi(ini, end)
7 total_pi = comm.allreduce(pi, op = MPI.SUM)

```

## 2.3 Provide runtime analysis on varying number of processes

### 2.3.1 The value of $\infty$ set as $10^2$

Processes	Time(s)
1	0,161
4	0,102
8	0,063
16	0,049
32	0,042

Table 2.1: Results for processes-time for the value of  $\infty$  set as  $10^2$




Figure 2.2: Combine results for the value of  $\infty$  set as  $10^2$

It can be seen in Table 2.2 how results improve with respect to time as the number of processes running in parallel increases.

### 2.3.2 The value of $\infty$ set as $10^3$

Processes	Time(s)
1	103,252
4	63,580
8	41,622
16	31,234
32	28,616

Table 2.2: Results for processes-time for the value of  $\infty$  set as  $10^3$




Figure 2.3: Combine results for the value of  $\infty$  set as  $10^3$

It can be seen in Table 2.3 how results improve with respect to time as the number of processes running in parallel increases.

### 2.3.3 Run program

It is possible to store all decimals of Pi in a file, since when it is indicated a very high number the console is break. This option is deactivated by default, to verify its operation with small values and to see them directly in console.

```
1 if rank== 0:
2     MPI.Finalize()
3     #     outfile = open('texto.txt', 'w')
4     #     outfile.write(str(total_pi))
5     #     print("Pi:\t", total_pi)
6     print("Sum processes time:\t{}\t{}".format(size, (totalTime/size)))
```

## 3 Matrix - matrix multiplication using collective communication

### 3.1 How you assign tasks to different processes?

To make a correct division of labor, the dimension of the matrix must be a **multiple** of the number of **processes**.

```
1 SIZE_ROWS = 3
2 ROWS = int(numpy.ceil(SIZE_ROWS/size)*size)
```

The division of labor will be done with respect to the columns of the first matrix (**matrix\_A**) and of processes size.

```
1 size = comm.Get_size()
2 n_calc = int(ROWS/size)
```

Matrix to scatter communication (matrix\_A chunks)

```
1 chunk_A = numpy.zeros((n_calc, COLS), dtype='i')
```

Broadcast of second matrix (matrix\_B, the same data to all processes) because only matrix\_A will be partitioned.

```
1 comm.Bcast([matrix_B, MPI.INT])
```

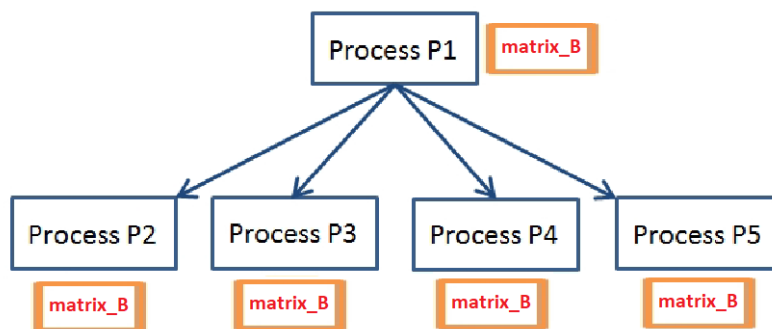


Figure 3.1: Broadcast of second matrix

Scatter of matrix\_A (matrix\_A chunks of data to each process)

```
1 comm.Scatter([matrix_A, MPI.INT], [chunk_A, MPI.INT])
```

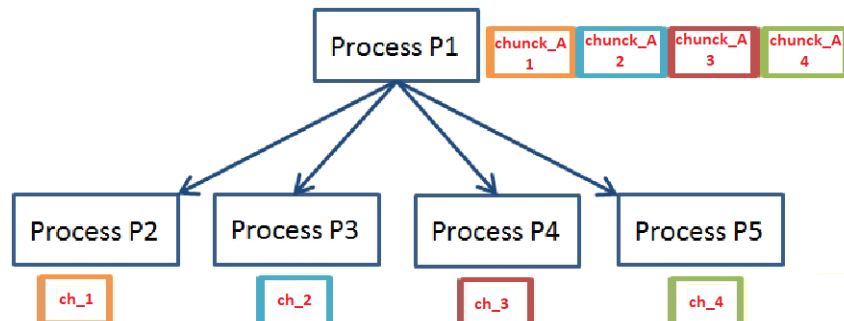


Figure 3.2: Scatter of first matrix

## 3.2 How you combine results from all the processes?

Once all the work has been distributed, each process performs its part of the multiplication, saving the result in an auxiliary matrix.

```
1 chunk_C = numpy.dot(chunk_A,matrix_B)
```

To combine all the results, I've used **Gather** (Figure 3.3). All processes send their auxiliary matrix to a root process that collects the data received and save it in the result matrix (**matrix\_C**).

```
1 comm.Gather([chunk_C,MPI.INT],[matrix_C,MPI.INT])
```

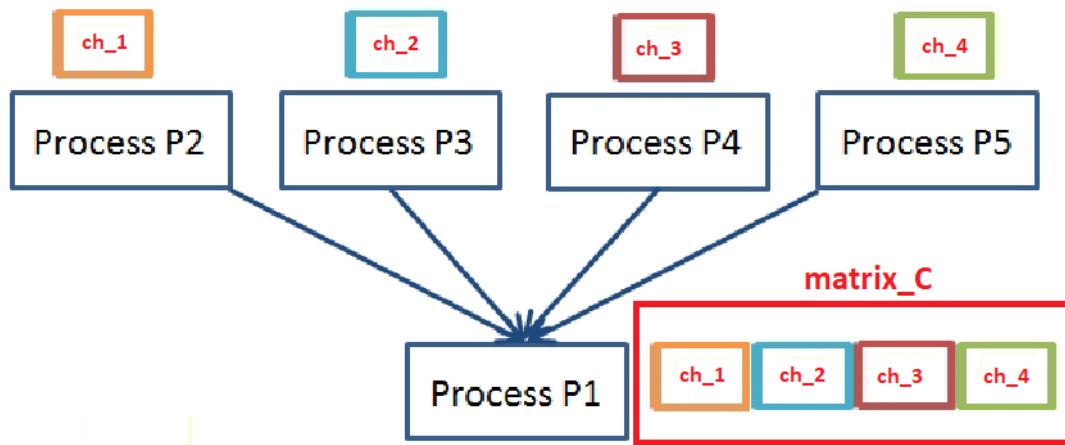


Figure 3.3: Gather of result matrix

## 3.3 Provide runtime analysis on varying number of processes

### 3.3.1 First run

$matrix\_A = 960 * 960$

$matrix\_B = 960 * 960$

$matrix\_A * B = 960 * 960 = matrix\_C$

Processes	Time(seconds)
1	13,582
4	6,029
8	5,678
16	4,868
32	4,808

Table 3.1: First run results time per processes



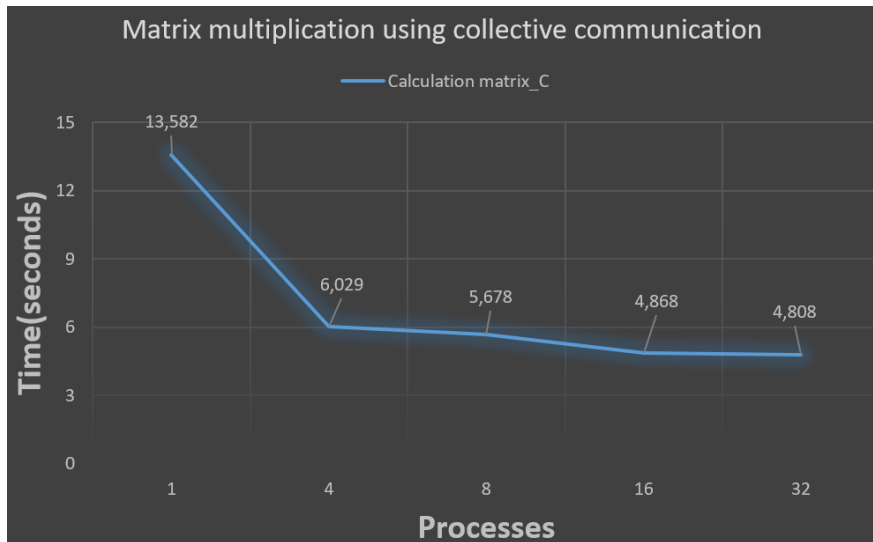


Figure 3.4: Combine first run results

### 3.3.2 Second run

$matrix\_A = 1920 * 960$

$matrix\_B = 960 * 1920$

$matrix\_A * B = 1920 * 1920 = matrix\_C$

Processes	Time(seconds)
1	58,746
2	34,593
4	25,280
16	21,060
32	20,809

Table 3.2: Second run results time per processes




Figure 3.5: Combine second run results

### 3.3.3 Third run

$matrix\_A = 3840 * 1920$

$matrix\_B = 1920 * 3840$

$matrix\_A * B = 3840 * 3840 = matrix\_C$

Processes	Time(seconds)
1	569,995
2	352,940
4	263,712
16	197,030
32	194,124

Table 3.3: Third run results time per processes




Figure 3.6: Combine third run results

### 3.3.4 Conclusions

It can be seen in Figure 3.6, Figure 3.5 and Figure 3.4 how results improve with respect to time as the number of processes running in parallel increases.

In all three cases it is coincided that the two contiguous points in which the greatest time difference is greatest is from one process to two processes.

On the other hand, also in the three cases it is coincided that the smallest decrease of time occurs from sixteen processes to thirty-two processes.