

Exercise Sheet 6

Mario Rodríguez Ruiz

May 26, 2017

Contents

1	Explain your system	3
2	Data cleaning and text tokenization	3
2.1	Cleaning: remove all punctuations and numbers	3
2.2	Stopping: removing meaningless words	3
2.2.1	mapper.py	4
2.2.2	reduce.py	6
3	Calculate TFIDF scores of words/tokens	6
3.1	Mapper step-by-step	6
3.1.1	mapper.py	7
3.2	Reducer step-by-step	8
3.2.1	reduce.py	10
3.3	Execution test	12

List of Figures

3.1	Eclipse run configurations	12
3.2	Hadoop input files	12
3.3	Results of the TFIDF calculation	13
3.4	Results of the TFIDF calculation tutorial	13

List of Tables

1.1	My system	3
-----	---------------------	---

1 Explain your system

Home System	
Machine	Asus Notebook ROG G60Jx
Operating System	Windows 10 Pro 64-bit
CPU	Intel Core i7 720QM @1.60GHz
Number of cores	4
Number of threads	8
RAM	16GB @665MHz (9-9-9-24)
Programming language version Python	v3.6.1:69c0db5 64 bit
Programming language version Java	v1.8
Hadoop	v2.6.0

Table 1.1: My system

2 Data cleaning and text tokenization

2.1 Cleaning: remove all punctuations and numbers

First to remove all punctuations and numbers, I've done this task in the Mapper. After reading each line the *sub* function is called, imported from *re* in which it specifies which characters are the only ones that can be included in each phrase.

```
1 import re
2
3 # input comes from STDIN
4 for line in sys.stdin:
5     # remove anything other than an alphabet letter
6     line = re.sub('[^A-Za-z]+', ' ', line)
7     [...]
```

To indicate which characters are allowed, the expression `'[A-Za-z]+'` is used. This indicates that everything that is not a letter of the alphabet is excluded.

```
1 line = re.sub('[^A-Za-z]+', ' ', line)
```

The following is to indicate by which character to replace each letter not allowed, in this case will be done with a space (' ') and finally the complete phrase to be examined (**line**).

2.2 Stopping: removing meaningless words

The most important thing at this point is knowing how to differentiate when it comes to one input and when another. That is, when it is processing the entry containing the Common English Words and when it is processing the file containing the data it want to count.

```
1 # split the line into words
2 words = line.split()
```

```

3 # storage line for english words
4 en_words = line
5
6 # if line size is high than 550 it is because
7 # the text file is the common english words
8 if(len(line)>550):
9     print ('{}\t{}'.format(en_words, 0))
10 # it is the text file to evaluate
11 else:
12     # each word found in the data file
13     for word in words:
14         if len(word)>1:
15             # write the results to STDOUT
16             print ('{}\t{}'.format(word, 1))

```

To differentiate this, what I have done has been to distinguish them through their size. As you know the minimum size (as columns) that has the Common Words file, it could be characterize the IF-switch for it.

As it could be can see in the code fragment above, if the size of Line is greater than 550 (columns) is because it is the file of Common Words. With this it can be already differentiated between inputs from the Mapper and then work comfortably on the Reducer.

2.2.1 mapper.py

```

1 import sys
2 import re
3
4 # input comes from STDIN
5 for line in sys.stdin:
6     # remove anything other than an alphabet letter
7     line = re.sub('[^A-Za-z]+', ' ', line)
8     # converts the text to lowercase
9     line = line.lower()
10    # split the line into words
11    words = line.split()
12    # storage line for english words
13    en_words = line
14
15    # if line size is high than 550 it is because
16    # the text file is the common english words
17    if(len(line)>550):
18        print ('{}\t{}'.format(en_words, 0))
19    # it is the text file to evaluate
20    else:
21        # each word found in the data file
22        for word in words:

```

```

23     if len(word)>1:
24         # write the results to STDOUT
25         print ('{}\t{}'.format(word, 1))

```

In the part of Reducer, there are two details to take into account: The first one is to know when the input from the Mapper is Common English Words or when it comes to the text file to be examined.

The second one is the way each word is counted.

```

1  # input comes from STDIN
2  for line in sys.stdin:
3      line = line.strip()
4      word, count = line.split('\t', 1)
5      # Converts input to numeric value
6      count = int(count)
7      # If count is zero it means that the input is the Common
      English
8      if count == 0:
9          [...]
10     # if count is one means that the entry is now the word bank
11     elif count == 1:
12         [...]

```

In order to solve the first question, a simple IF-switch has been used. Once the input from the Mapper in which both the text and a counter (or differentiator) have been stored, these results are analyzed.

If the second input is equal to zero, this variable works as a differentiator. That is, the entry is about the common words english.

On the contrary, if it is a number one it means that the file is the data file and that it will work as a counter.

```

1  # If count is zero it means that the input is the Common
    English
2  elif count == 1:
3      [...]
4      elif word not in common_english_words:
5          if current_word:
6              # write result to STDOUT
7              print ('{}\t{}'.format(current_word, current_count))
8              current_count = count
9              current_word = word

```

Once the Common Words have been stored in the Reducer, the check can be performed.

As seen in the previous code fragment, for each word it is checked if it is in the array of Common Words English. Only in the case of not being in these would store a new word in the Reducer.

2.2.2 reduce.py

```
1 import sys
2
3 current_word = None
4 current_count = 0
5 word = None
6 common_english_words = []
7
8 # input comes from STDIN
9 for line in sys.stdin:
10     line = line.strip()
11     word, count = line.split('\t', 1)
12     # Converts input to numeric value
13     count = int(count)
14     # If count is zero it means that the input is the Common
15     # English
16     if count == 0:
17         common_english_words = word.split(' ')
18     # if count is one means that the entry is now the word bank
19     elif count == 1:
20         # it works because Hadoop sorts map output
21         # by key (word) before it's passed to the reducer
22         if current_word == word:
23             current_count += count
24         # If the word isn't a Common English
25         elif word not in common_english_words:
26             if current_word:
27                 # write result to STDOUT
28                 print ('{}\t{}'.format(current_word, current_count))
29                 current_count = count
30                 current_word = word
31
32 # write the last word if needed
33 if current_word == word:
34     print ('{}\t{}'.format(current_word, current_count))
```

3 Calculate TFIDF scores of words/tokens

3.1 Mapper step-by-step

The first part of the **mapper** is the same as in the previous exercise. Where it starts to differentiate is in the part where the name of the input file is obtained.

This name is necessary to be able to classify the punctuation of the words according to the documents which contains them and, in addition, have counted the documents that have been processed so that can then apply the mathematical formulas correctly.

```
1 # Find the name of the input file
2 try:
3     input_file = os.environ['mapreduce_map_input_file']
4 except KeyError:
5     input_file = os.environ['map_input_file']
```

The rest of the code is similar to the one presented in the previous exercise with the exception of a single detail: now the **input processed file name** is also passed in the output in order to work with it in the reducer.

```
1 # write the results to STDOUT
2 print ('{}\t{}\t{}'.format(word, 1, input_file))
```

It should be remembered that **one or a zero** was also write in the output (in addition to each word) depending on the file: a zero when it was the CEW file and a one when it was a normal file to process.

3.1.1 mapper.py

```
1 import sys
2 import re
3 import os
4
5 # input comes from STDIN
6 for line in sys.stdin:
7     # remove anything other than an alphabet letter
8     line = re.sub('[^A-Za-z]+', ' ', line)
9     # converts the text to lowercase
10    line = line.lower()
11    # split the line into words
12    words = line.split()
13    # storage line for english words
14    en_words = line
15
16    # Find the name of the input file
17    try:
18        input_file = os.environ['mapreduce_map_input_file']
19    except KeyError:
20        input_file = os.environ['map_input_file']
21
22    # if line size is high than 550 it is because
23    # the text file is the common english words
24    if(len(line)>550):
```

```

25     # write the results to STDOUT
26     print ('{}\t{}\t{}'.format(en_words, 0, input_file))
27 else:
28     # For each word of the line
29     for word in words:
30         if len(word)>1:
31             # write the results to STDOUT
32             print ('{}\t{}\t{}'.format(word, 1, input_file))

```

3.2 Reducer step-by-step

To perform this task, it has been convenient to use some shared variables to be able to collect all the necessary data.

```

1 # To save the CEW list
2 common_english_words = []
3 # To save the contents of each of the documents
4 docs = {}
5 # To save the input file names and use them as indexes
6 inputs = []
7 # number of inputs documents (without accounting for the CEW)
8 n_doc = 0

```

The fundamental first is "**common_english_words**", in which will be stored the words that doesn't want to count; the second "**docs**" to store the contents of each of the documents; the third "**inputs**" to store each of the names of the input files to use as indexes of the lists; and finally "**n_doc**" that counts the number of documents analyzed.

When adding a new input in the reducer (input_) he has had to specify in its corresponding line. Changes to the previous version (WordCount) when processing the entry are as follows: The first thing to check is if the file is the CEW (through the variable "count" explained above).

```

1 wordd, count, input_ = line.split('\t')
2 # Converts input to numeric value
3 count = int(count)
4
5 # If count is zero it means that the input is the Common
   English
6 if count == 0:
7     common_english_words = wordd.split(' ')
8 # If count is one and the name of the input file isn't yet
   known
9 elif count == 1 and input_ not in inputs and wordd not in
   common_english_words:
10     # Save the new name in the list of names
11     inputs.append(input_)

```



```

12 # Initialize new list with index of the new name
13 docs[input_] = wordd + ' '
14 # If count is one and the word isn't a CEW
15 elif count == 1 and wordd not in common_english_words:
16     # Add new word to the end of the list identified through the
        index
17 docs[input_] += wordd + ' '

```

If it is not the file, it is checked by its own name if it has already been analyzed. If it is not already known, a new list is initialized with its name as index and the first word that is processed.

Always keep in mind before adding a new word other than a CEW.

Once all the inputs have been processed, all that is left is to perform the calculations.

```

1 bloblist = []
2 # Stores the inputs documents (no CEW) in TextBlob format
3 for i in inputs:
4     i = tb(docs[i])
5     bloblist.append(i)

```

To do this, the lists that contain the words of each document will be transformed into **TextBlob format** to perform operations in a more manageable way through this library.

Finally, the function tfidf is called, in which operations are performed for each word based on all the documents that have been processed in the input (except the CEW).

```

1 # For each word of each document it calculates the TFIDF score
2 for i, blob in enumerate(bloblist):
3     print("\nTop words in document {}".format(inputs[i]))
4     scores = {word: tfidf(word, blob, bloblist) for word in blob
        .words}
5     # sort from highest to lowest list by tfidf scores
6     scores = sorted(scores.items(), key=lambda x: x[1], reverse=
        True)
7     print("\t-----")
8     print("\tWord\tTF-IDF")
9     print("\t-----")
10    for word, score in scores[:3]:
11        print("\t{}\t{}".format(word, round(score, 3)))

```

Term frequency which is the number of times a word appears in a document:

```

1 def tf(word, blob):
2     return blob.words.count(word) / len(blob.words)

```

Number of documents containing a specific word

```
1 def n_containing(word, bloblist):
2     return sum(1 for blob in bloblist if word in blob.words)
```

Inverse document frequency which measures how common a word is among all processed documents:

```
1 def idf(word, bloblist):
2     return math.log(len(bloblist) / (1 + n_containing(word,
    bloblist)))
```

Computes the TF-IDF score:

```
1 def tfidf(word, blob, bloblist):
2     return tf(word, blob) * idf(word, bloblist)
```

3.2.1 reduce.py

```
1 import sys
2 import math
3 from textblob import TextBlob as tb
4
5 # To save the CEW list
6 common_english_words = []
7 # To save the contents of each of the documents
8 docs = {}
9 # To save the input file names and use them as indexes
10 inputs = []
11 # number of inputs documents (without accounting for the CEW)
12 n_doc = 0
13
14 # term frequency which is the number of times a word appears
    in a document blob
15 def tf(word, blob):
16     return blob.words.count(word) / len(blob.words)
17
18 # number of documents containing word
19 def n_containing(word, bloblist):
20     return sum(1 for blob in bloblist if word in blob.words)
21
22 # inverse document frequency which measures how common a word
    is among all documents in bloblist
23 def idf(word, bloblist):
24     return math.log(len(bloblist) / (1 + n_containing(word,
    bloblist)))
25
26 # computes the TF-IDF score
```

```

27 def tfidf(word, blob, bloblist):
28     return tf(word, blob) * idf(word, bloblist)
29
30 # input comes from STDIN
31 for line in sys.stdin:
32     line = line.strip()
33     wordd, count, input_ = line.split('\t')
34     # Converts input to numeric value
35     count = int(count)
36
37     # If count is zero it means that the input is the Common
        English
38     if count == 0:
39         common_english_words = wordd.split(' ')
40     # If count is one and the name of the input file isn't yet
        known
41     elif count == 1 and input_ not in inputs and wordd not in
        common_english_words:
42         # Save the new name in the list of names
43         inputs.append(input_)
44         # Initialize new list with index of the new name
45         docs[input_] = wordd + ' '
46     # If count is one and the word isn't a CEW
47     elif count == 1 and wordd not in common_english_words:
48         # Add new word to the end of the list identified through
        the index
49         docs[input_] += wordd + ' '
50
51 bloblist = []
52 # Stores the inputs documents (no CEW) in TextBlob format
53 for i in inputs:
54     i = tb(docs[i])
55     bloblist.append(i)
56
57 # For each word of each document it calculates the TFIDF score
58 for i, blob in enumerate(bloblist):
59     print("\nTop words in document {}".format(inputs[i]))
60     scores = {word: tfidf(word, blob, bloblist) for word in blob
        .words}
61     # sort from highest to lowest list by tfidf scores
62     scores = sorted(scores.items(), key=lambda x: x[1], reverse=
        True)
63     print("\t-----")
64     print("\tWord\tTF-IDF")
65     print("\t-----")
66     for word, score in scores[:3]:
67         print("\t{}\t{}".format(word, round(score, 3)))

```

3.3 Execution test

Image 3.1 shows the execution command that has been carried out through Eclipse.

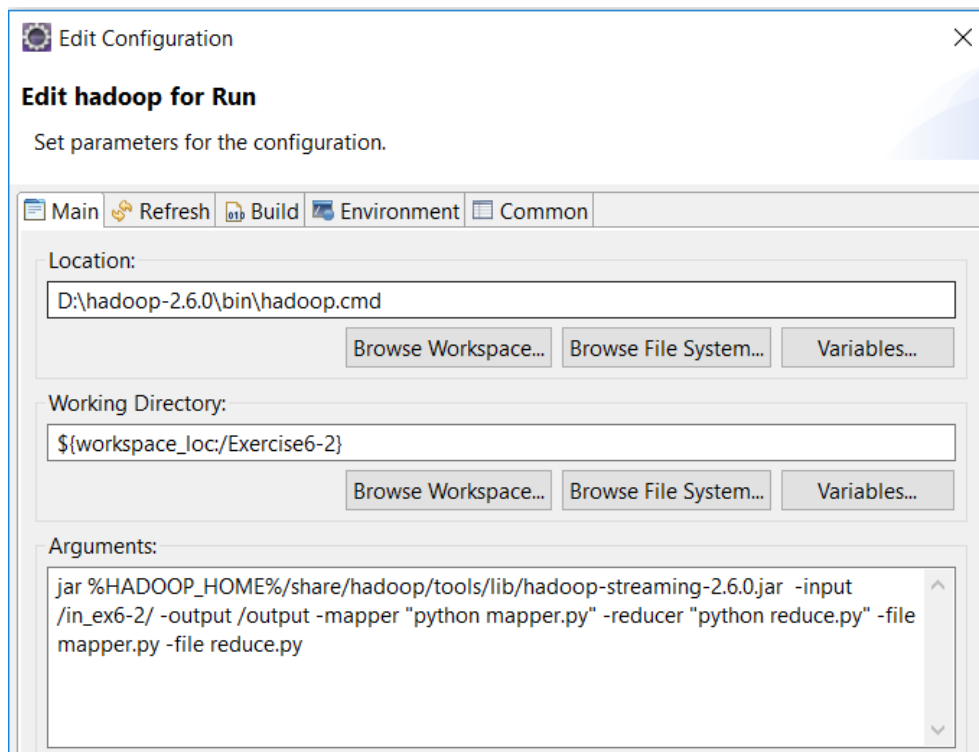


Figure 3.1: Eclipse run configurations

Image 3.2 shows the input files to be processed.

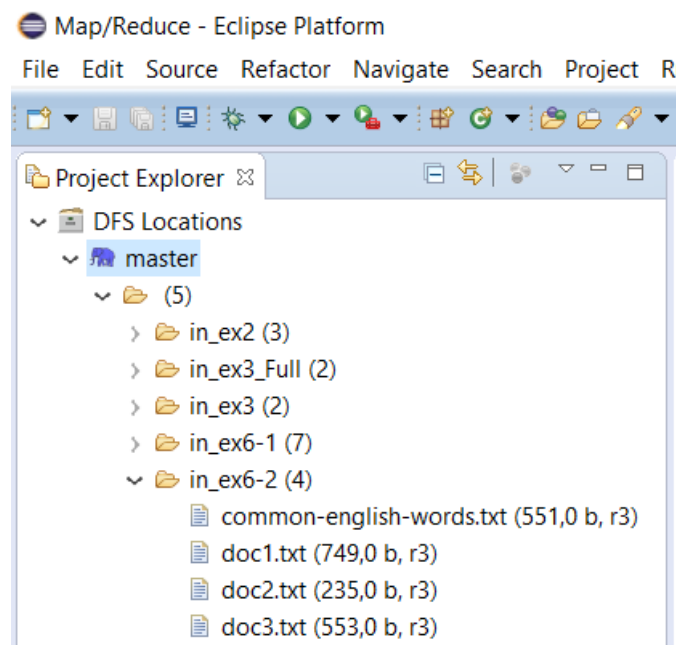


Figure 3.2: Hadoop input files

Image 3.3 shows the output files with the results of the TFIDF calculation.

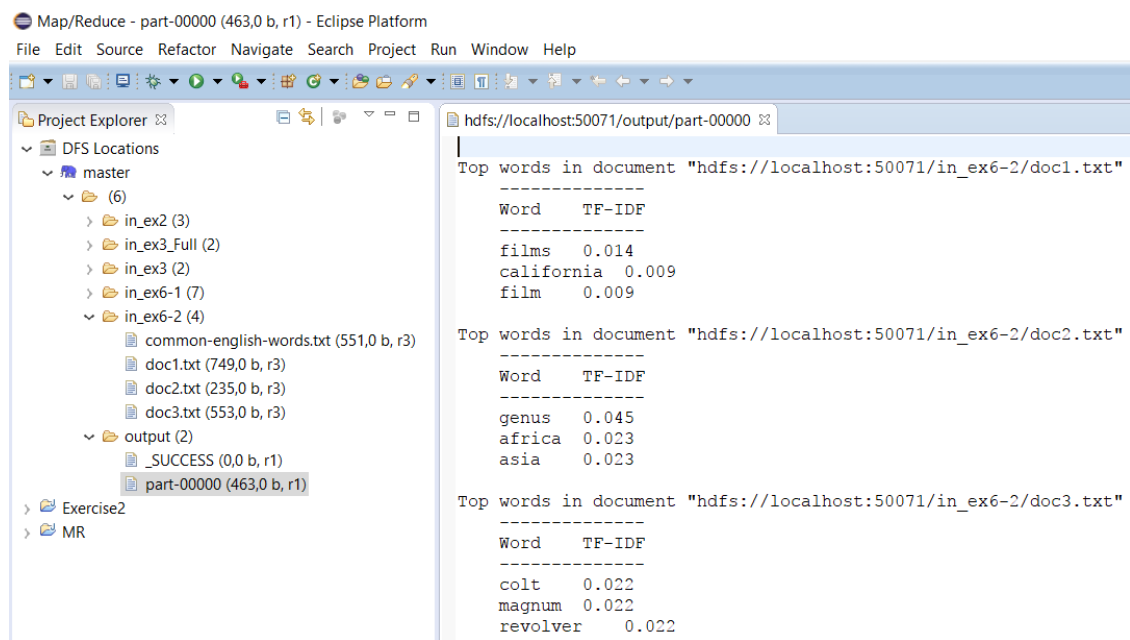


Figure 3.3: Results of the TFIDF calculation

Image 3.4 shows the results of the TFIDF calculation of tutorial <http://stevenloria.com/finding-important-words-in-a-document-using-tf-idf/>

```
Top words in document 1
Word: films, TF-IDF: 0.00997
Word: film, TF-IDF: 0.00665
Word: California, TF-IDF: 0.00665

Top words in document 2
Word: genus, TF-IDF: 0.02192
Word: among, TF-IDF: 0.01096
Word: Currently, TF-IDF: 0.01096

Top words in document 3
Word: Magnum, TF-IDF: 0.01382
Word: revolver, TF-IDF: 0.01382
Word: Colt, TF-IDF: 0.01382
```

Figure 3.4: Results of the TFIDF calculation tutorial

The results are not exactly the same because this tutorial takes into account all the characters, being also the words of the **common english word**.